CSE 488 (Section 1)
[Summer 2022]

Lab 02 Assignment Submission Report

Assignment Title: MapReduce Programming

Submitted by:
Rizvee Hassan Prito
2019-3-60-041

1. Screenshots

Problem 1:

Driver class:

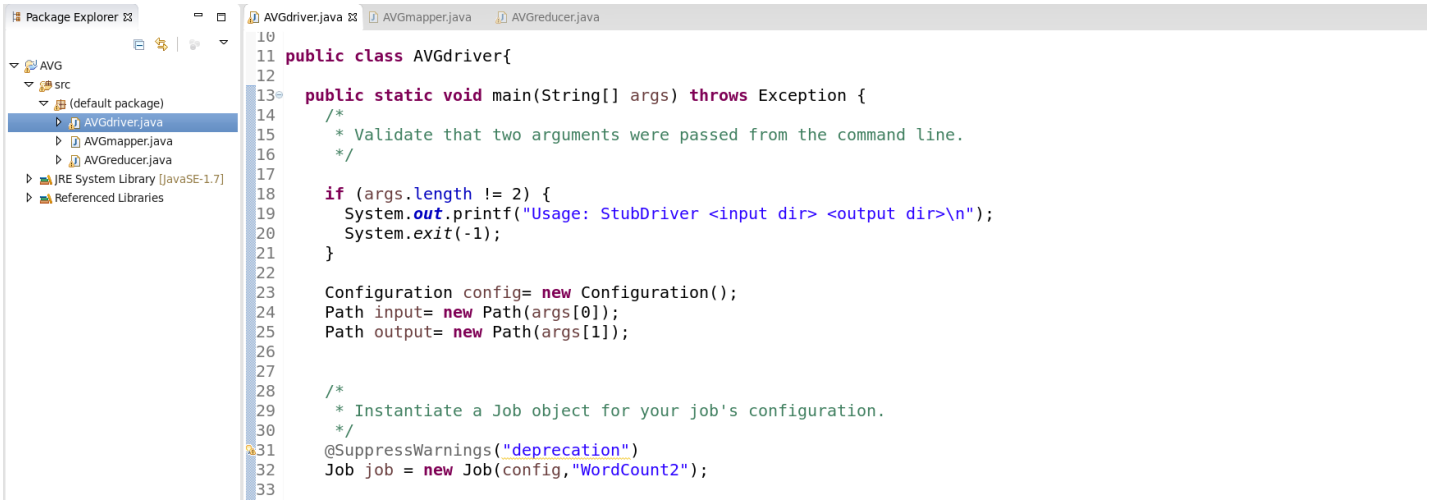

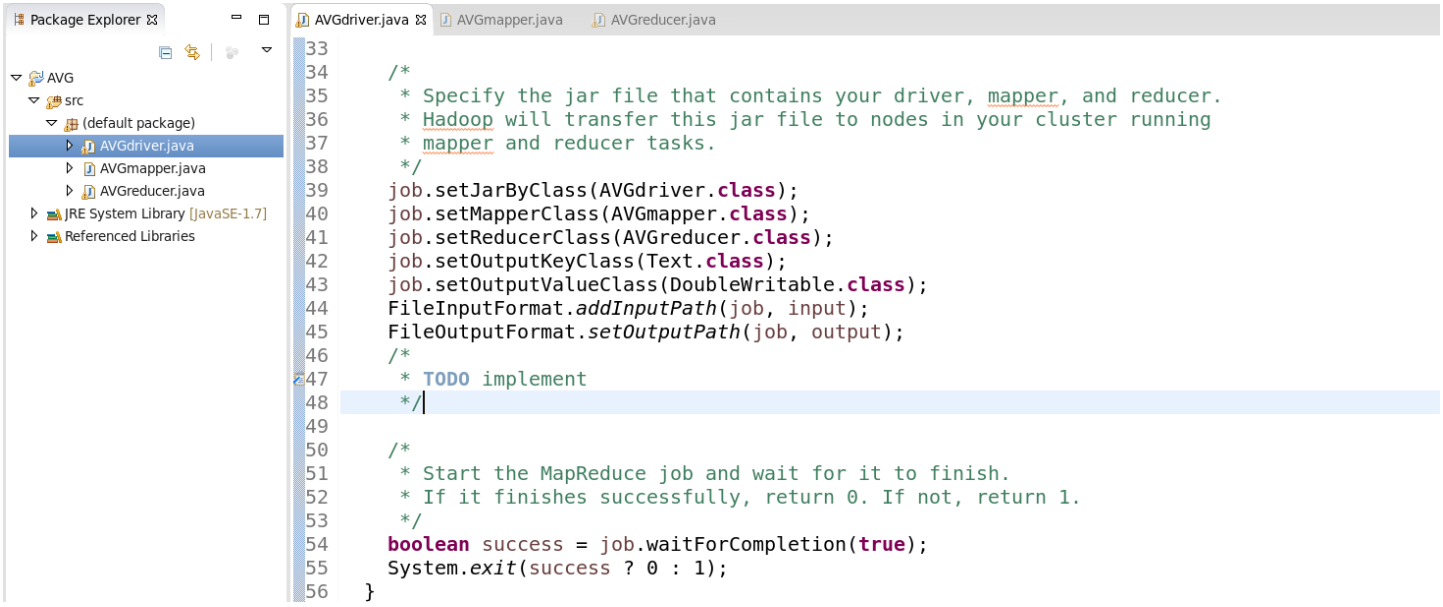

Mapper class:

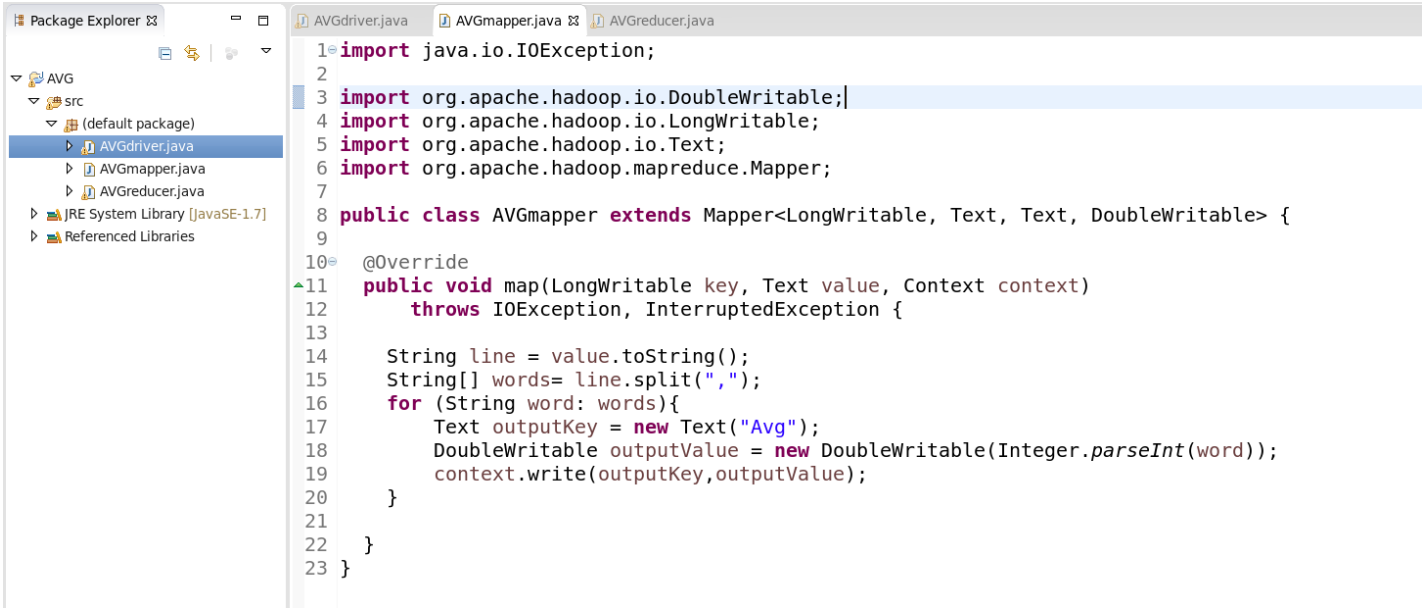

Reducer class:

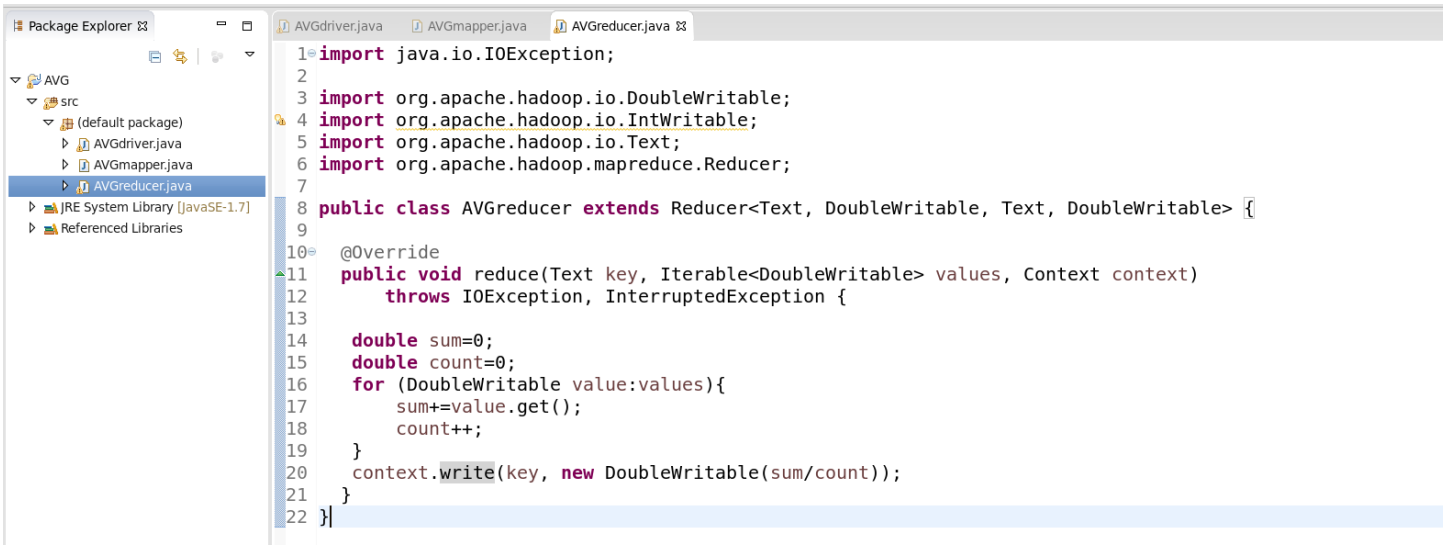

Input File:

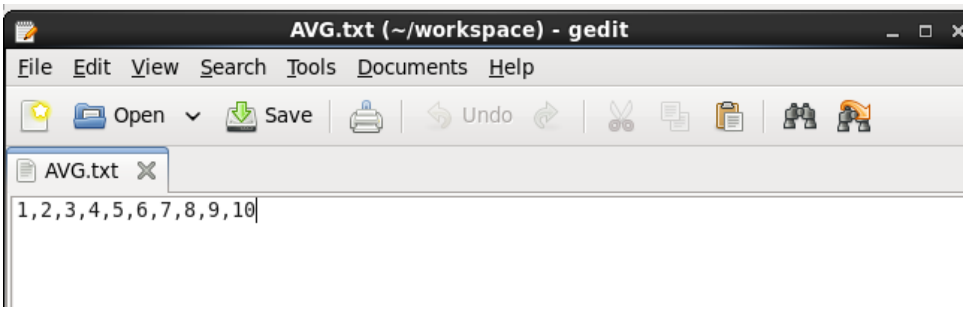

Terminal:

```
22/06/26 05:25:42 INFO mapreduce.Job:  map 0% reduce 0%
22/06/26 05:25:51 INFO mapreduce.Job:  map 100% reduce 0%
22/06/26 05:26:02 INFO mapreduce.Job:  map 100% reduce 100%
22/06/26 05:26:03 INFO mapreduce.Job: Job job_1656243722115_0002 completed successfully
22/06/26 05:26:04 INFO mapreduce.Job: Counters: 49
        File System Counters
                FILE: Number of bytes read=146
                FILE: Number of bytes written=223019
                FILE: Number of read operations=0
                FILE: Number of large read operations=0
                FILE: Number of write operations=0
                HDFS: Number of bytes read=139
                HDFS: Number of bytes written=8
                HDFS: Number of read operations=6
                HDFS: Number of large read operations=0
                HDFS: Number of write operations=2
        Job Counters
                Launched map tasks=1
                Launched reduce tasks=1
                Data-local map tasks=1
                Total time spent by all maps in occupied slots (ms)=6732
                Total time spent by all reduces in occupied slots (ms)=8178
                Total time spent by all map tasks (ms)=6732
                Total time spent by all reduce tasks (ms)=8178
                Total vcore-seconds taken by all map tasks=6732
                Total vcore-seconds taken by all reduce tasks=8178
                Total megabyte-seconds taken by all map tasks=6893568
                Total megabyte-seconds taken by all reduce tasks=8374272
                Total megabyte-seconds taken by all reduce tasks=8374272
        Map-Reduce Framework
                Map input records=1
                Map output records=10
                Map output bytes=120
                Map output materialized bytes=146
                Input split bytes=118
                Combine input records=0
                Combine output records=0
                Reduce input groups=1
                Reduce shuffle bytes=146
                Reduce input records=10
                Reduce output records=1
                Spilled Records=20
                Shuffled Maps =1
                Failed Shuffles=0
                Merged Map outputs=1
                GC time elapsed (ms)=176
                CPU time spent (ms)=1650
                Physical memory (bytes) snapshot=334729216
                Virtual memory (bytes) snapshot=3008466944
                Total committed heap usage (bytes)=226365440
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=21
        File Output Format Counters
                Bytes Written=8
```

Output:

Home / user / cloudera / AVG / **part-r-00000**

```
Avg     5.5
```

Problem 2:

Driver class:

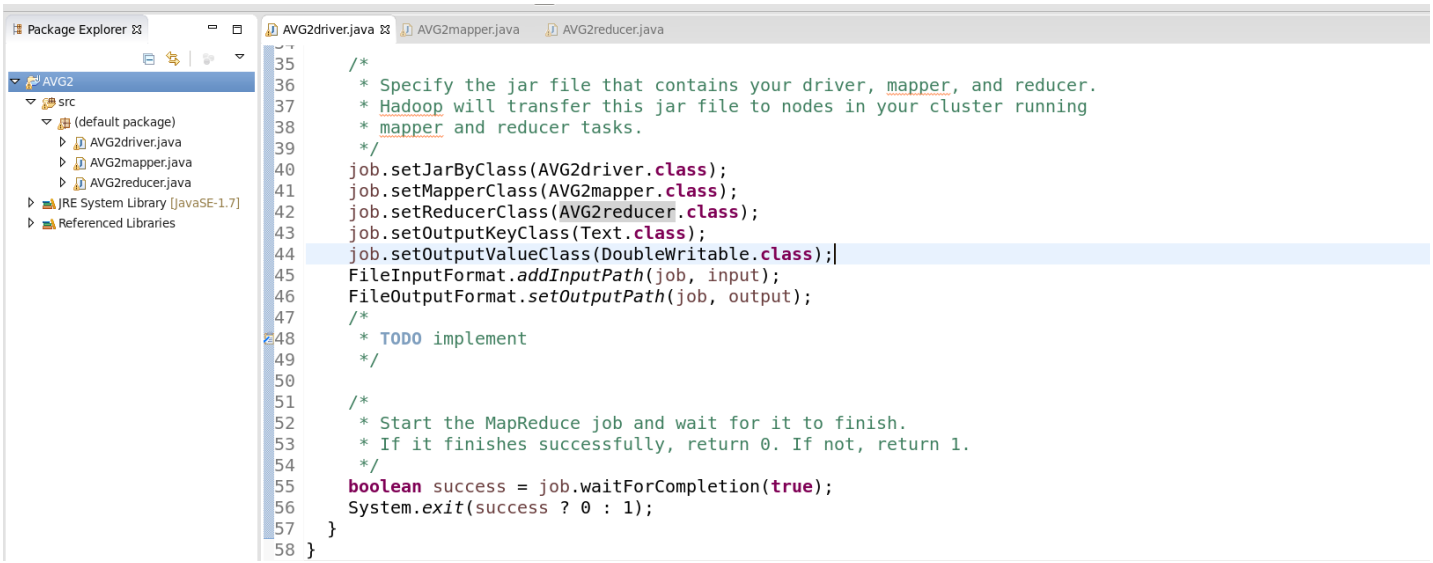

Mapper class:

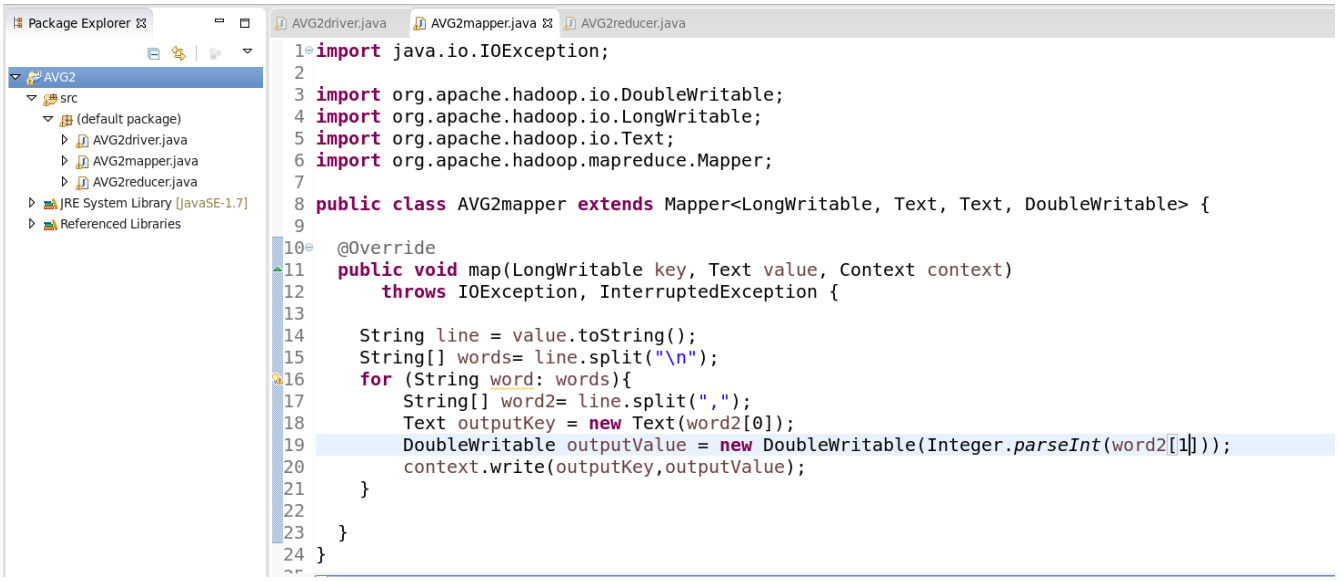

Reducer class:

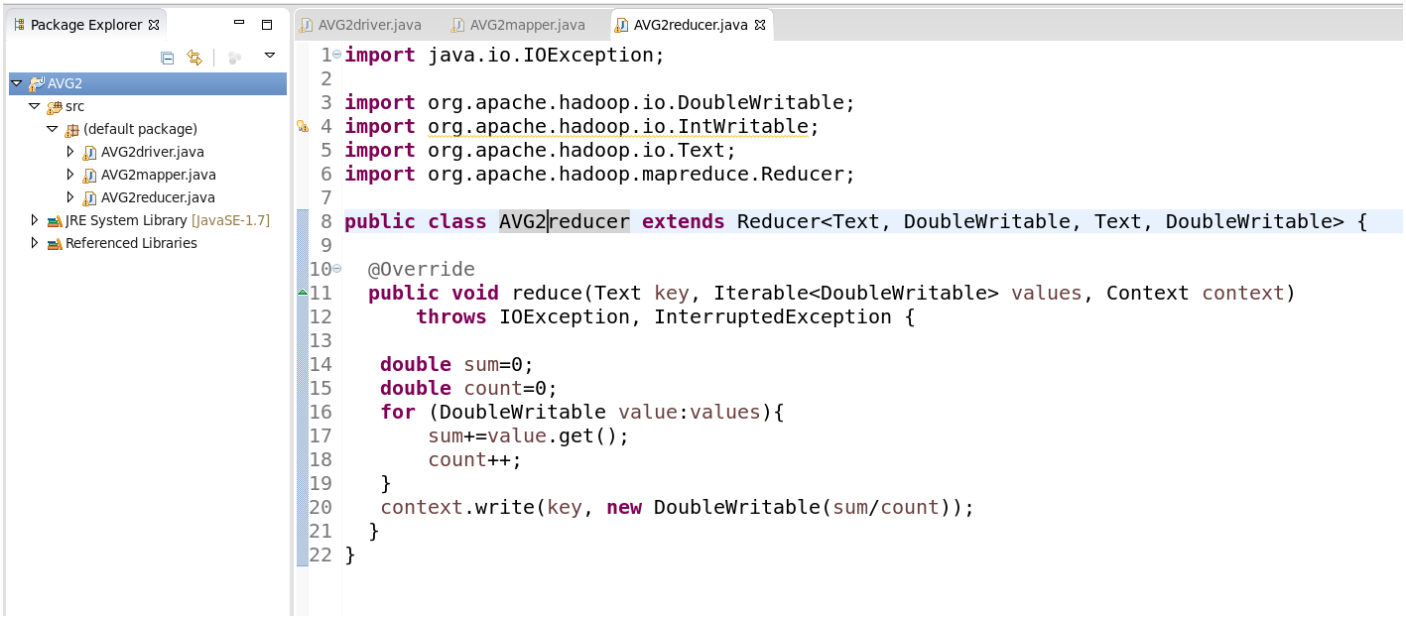

Input File:

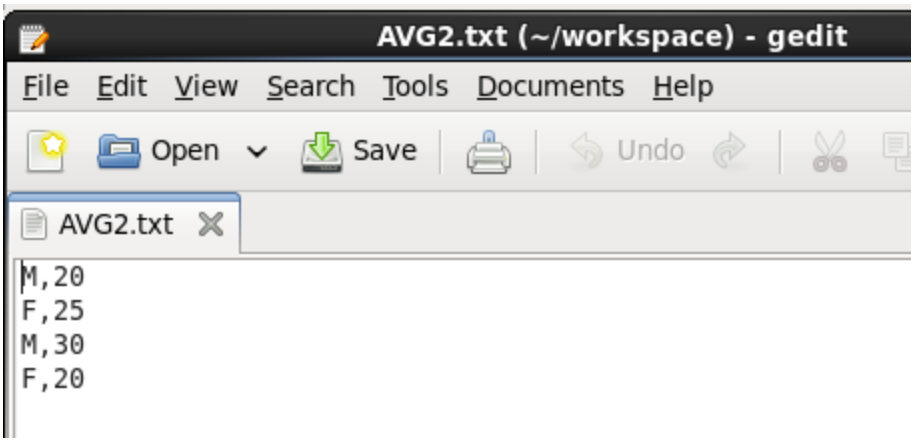

Terminal:

```
22/06/26 08:26:17 INFO mapreduce.Job: Job job_1656256202977_0001 running in uber mode : false
22/06/26 08:26:17 INFO mapreduce.Job:  map 0% reduce 0%
22/06/26 08:26:25 INFO mapreduce.Job:  map 100% reduce 0%
22/06/26 08:26:36 INFO mapreduce.Job:  map 100% reduce 100%
22/06/26 08:26:36 INFO mapreduce.Job: Job job_1656256202977_0001 completed successfully
22/06/26 08:26:36 INFO mapreduce.Job: Counters: 49
        File System Counters
                FILE: Number of bytes read=54
                FILE: Number of bytes written=222843
                FILE: Number of read operations=0
                FILE: Number of large read operations=0
                FILE: Number of write operations=0
                HDFS: Number of bytes read=139
                HDFS: Number of bytes written=14
                HDFS: Number of read operations=6
                HDFS: Number of large read operations=0
                HDFS: Number of write operations=2
        Job Counters
                Launched map tasks=1
                Launched reduce tasks=1
                Data-local map tasks=1
                Total time spent by all maps in occupied slots (ms)=6087
                Total time spent by all reduces in occupied slots (ms)=7816
                Total time spent by all map tasks (ms)=6087
                Total time spent by all reduce tasks (ms)=7816
                Total vcore-seconds taken by all map tasks=6087
                Total vcore-seconds taken by all reduce tasks=7816
                Total megabyte-seconds taken by all map tasks=6233088
                Total megabyte-seconds taken by all reduce tasks=8003584
```

```
Map-Reduce Framework
        Map input records=4
        Map output records=4
        Map output bytes=40
        Map output materialized bytes=54
        Input split bytes=119
        Combine input records=0
        Combine output records=0
        Reduce input groups=2
        Reduce shuffle bytes=54
        Reduce input records=4
        Reduce output records=2
        Spilled Records=8
        Shuffled Maps =1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=176
        CPU time spent (ms)=1530
        Physical memory (bytes) snapshot=347967488
        Virtual memory (bytes) snapshot=3008630784
        Total committed heap usage (bytes)=226365440
Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
File Input Format Counters
        Bytes Read=20
File Output Format Counters
        Bytes Written=14
```

Output:

Home / user / cloudera / AVG2 / **part-r-00000**

```
F       22.5
M       25.0
```

Problem 3:

Driver class:

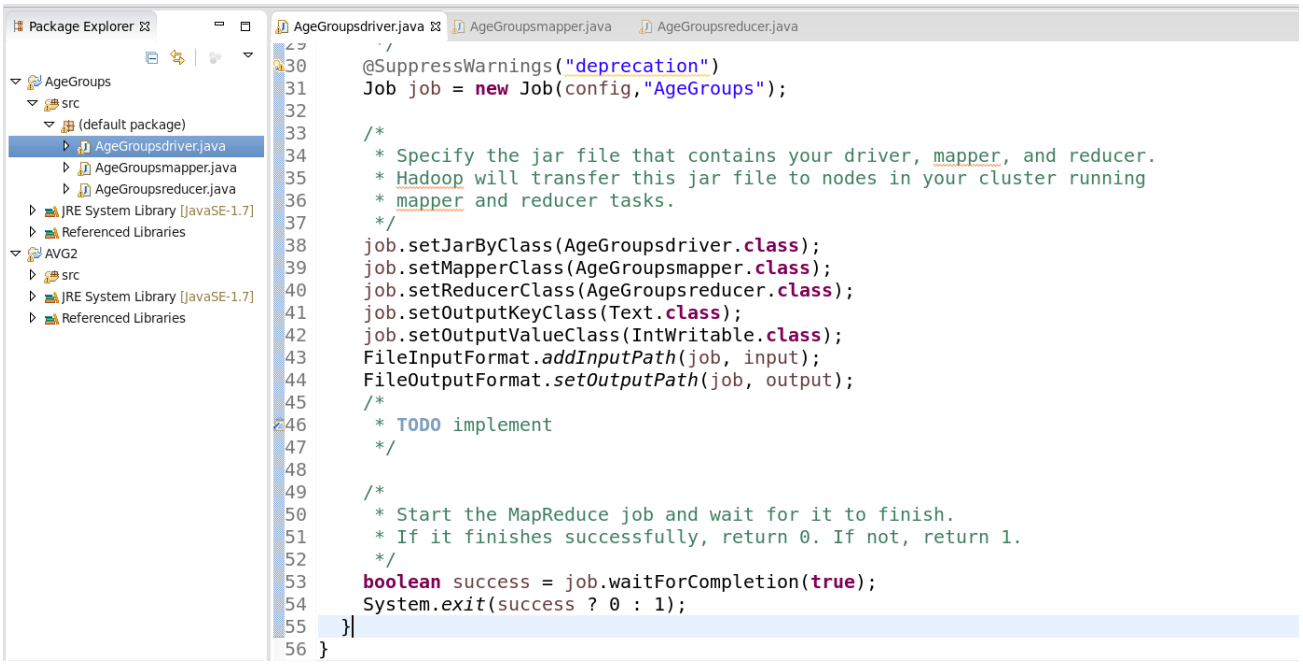

Mapper class:

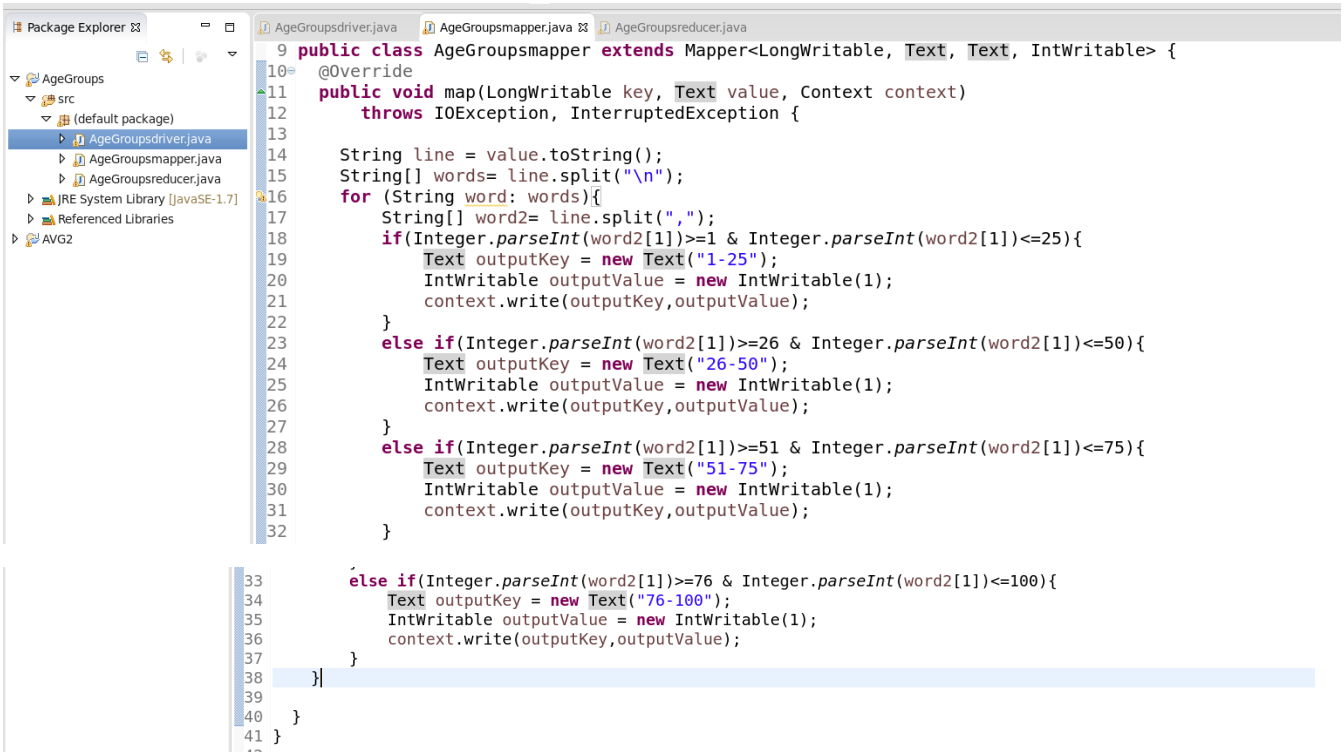

Reducer class:

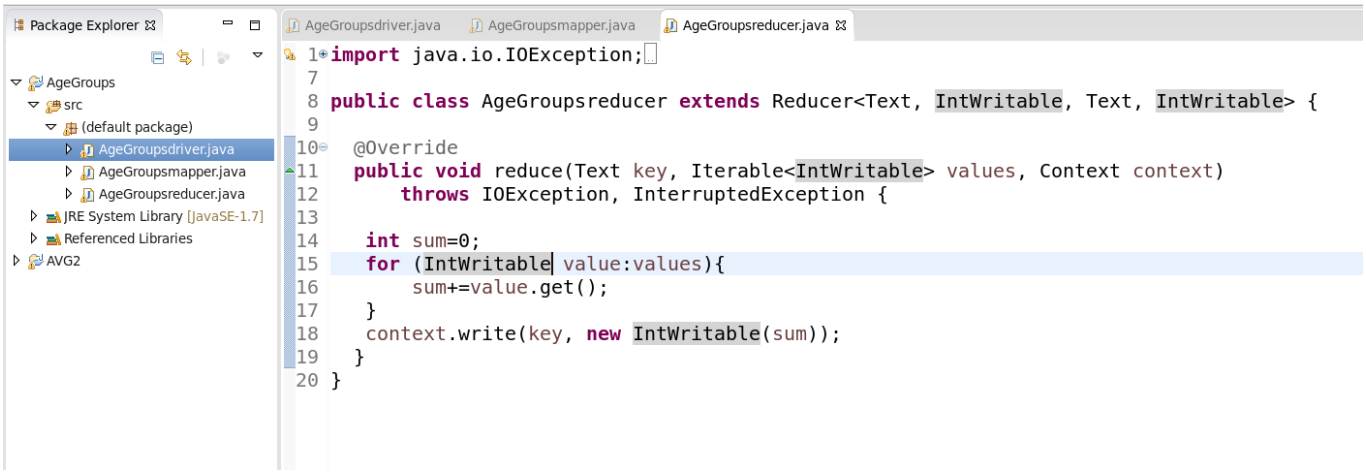

Input File:

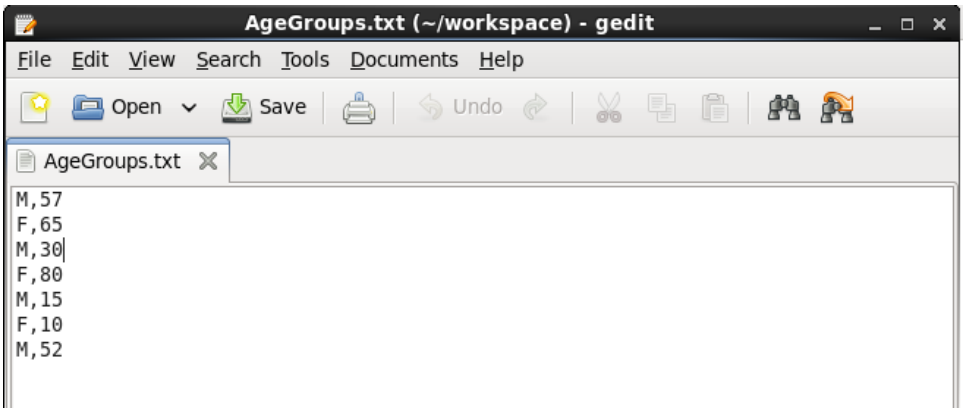

Terminal:

```
22/06/26 08:51:50 INFO mapreduce.Job:  map 0% reduce 0%
22/06/26 08:51:59 INFO mapreduce.Job:  map 100% reduce 0%
22/06/26 08:52:10 INFO mapreduce.Job:  map 100% reduce 100%
22/06/26 08:52:10 INFO mapreduce.Job: Job job_1656256202977_0003 completed successfully
22/06/26 08:52:11 INFO mapreduce.Job: Counters: 49
        File System Counters
                FILE: Number of bytes read=89
                FILE: Number of bytes written=222947
                FILE: Number of read operations=0
                FILE: Number of large read operations=0
                FILE: Number of write operations=0
                HDFS: Number of bytes read=159
                HDFS: Number of bytes written=32
                HDFS: Number of read operations=6
                HDFS: Number of large read operations=0
                HDFS: Number of write operations=2
        Job Counters
                Launched map tasks=1
                Launched reduce tasks=1
                Data-local map tasks=1
                Total time spent by all maps in occupied slots (ms)=6859
                Total time spent by all reduces in occupied slots (ms)=7459
                Total time spent by all map tasks (ms)=6859
                Total time spent by all reduce tasks (ms)=7459
                Total vcore-seconds taken by all map tasks=6859
                Total vcore-seconds taken by all reduce tasks=7459
                Total megabyte-seconds taken by all map tasks=7023616
                Total megabyte-seconds taken by all reduce tasks=7638016
        Map-Reduce Framework
                Map input records=7
                Map output records=7
                Map output bytes=69
                Map output materialized bytes=89
                Input split bytes=124
                Combine input records=0
                Combine output records=0
                Reduce input groups=4
                Reduce shuffle bytes=89
                Reduce input records=7
                Reduce output records=4
                Spilled Records=14
                Shuffled Maps =1
                Failed Shuffles=0
                Merged Map outputs=1
                GC time elapsed (ms)=182
                CPU time spent (ms)=1550
                Physical memory (bytes) snapshot=357978112
                Virtual memory (bytes) snapshot=3008475136
                Total committed heap usage (bytes)=226365440
        Shuffle Errors
                BAD_ID=0
                CONNECTION=0
                IO_ERROR=0
                WRONG_LENGTH=0
                WRONG_MAP=0
                WRONG_REDUCE=0
        File Input Format Counters
                Bytes Read=35
        File Output Format Counters
                Bytes Written=32
```

Output:

Home / user / cloudera / AgeGroups / **part-r-00000**

```
1-25    2
26-50   1
51-75   3
76-100  1
```

Learning Outcome:

From this lab we have learned how to solve problems by using MapReduce programming style and execute the program in Hadoop Framework.